

Local RAG Assistant FAQ: How do I add my own documents to the assistant?

Copy your files into the documents folder and run `python ingest.py`. Supported formats are `.txt`, `.md`, `.pdf` and `.docx`. Ingestion rebuilds the whole knowledge base, so removed files disappear from answers too.

Local RAG Assistant FAQ: Which models does the assistant use?

Chat answers come from phi-3.5-mini and embeddings from qwen3-embedding-0.6b. Both run fully offline through Microsoft Foundry Local, which resolves each alias to the best variant for the local hardware (CPU, GPU or NPU).

Local RAG Assistant FAQ: Why does the assistant sometimes refuse to answer?

Two safety layers cause refusals. First, chunks whose cosine similarity to the question is below 0.30 are discarded; if nothing remains, a fixed fallback answer is returned without calling the model. Second, the system prompt instructs the model to say it does not know when the retrieved context does not contain the answer.

Local RAG Assistant FAQ: How can I make answers faster?

To speed up responses, retrieve fewer chunks by lowering TOP_K, reduce max_tokens so answers are shorter, or switch CHAT_MODEL_ALIAS to a smaller chat model such as qwen3-1.7b. Most latency comes from token generation, not retrieval.